

# CS107, Lecture 3

## Bits and Bytes; Bitwise Operators

reading:

*Bryant & O'Hallaron, Ch. 2.1*

# Plan For Today

- **Recap:** Integer Representations
- Truncating and Expanding
- Bitwise Boolean Operators and Masks
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- Bit Shift Operators

# Plan For Today

- **Recap:** Integer Representations
  - Truncating and Expanding
  - Bitwise Boolean Operators and Masks
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- Bit Shift Operators

# Base 2

1 0 1 1  
 $2^3$      $2^2$      $2^1$      $2^0$

# Hexadecimal

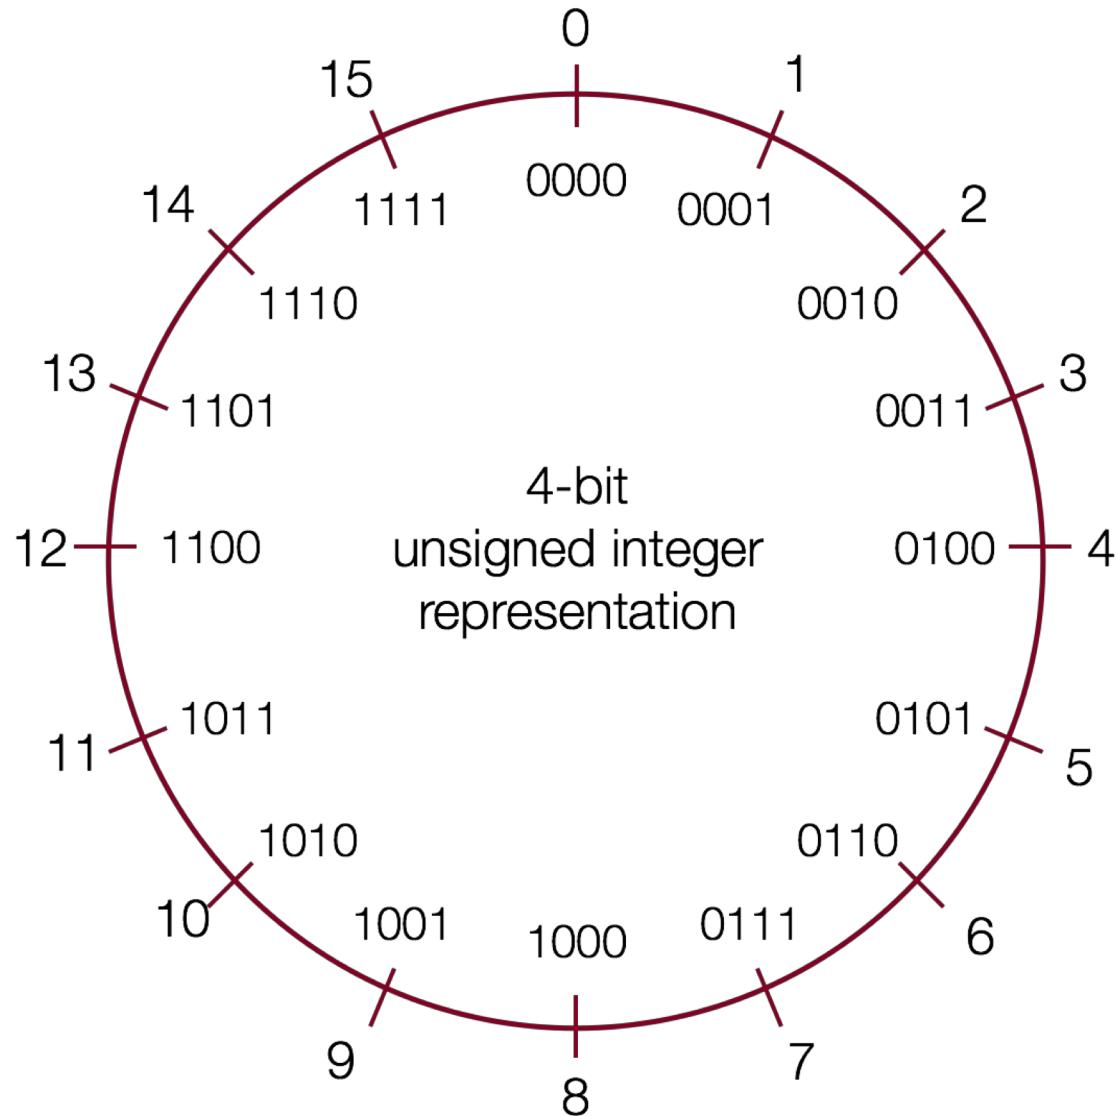
Hex digit	0	1	2	3	4	5	6	7
Decimal value	0	1	2	3	4	5	6	7
Binary value	0000	0001	0010	0011	0100	0101	0110	0111
Hex digit	8	9	A	B	C	D	E	F
Decimal value	8	9	10	11	12	13	14	15
Binary value	1000	1001	1010	1011	1100	1101	1110	1111

# Number Representations

C declaration			Bytes	
Signed	Unsigned	32-bit	64-bit	
[signed] char	unsigned char	1	1	
short	unsigned short	2	2	
int	unsigned	4	4	
long	unsigned long	4	8	
int32_t	uint32_t	4	4	
int64_t	uint64_t	8	8	
char *		4	8	
float		4	4	
double		8	8	

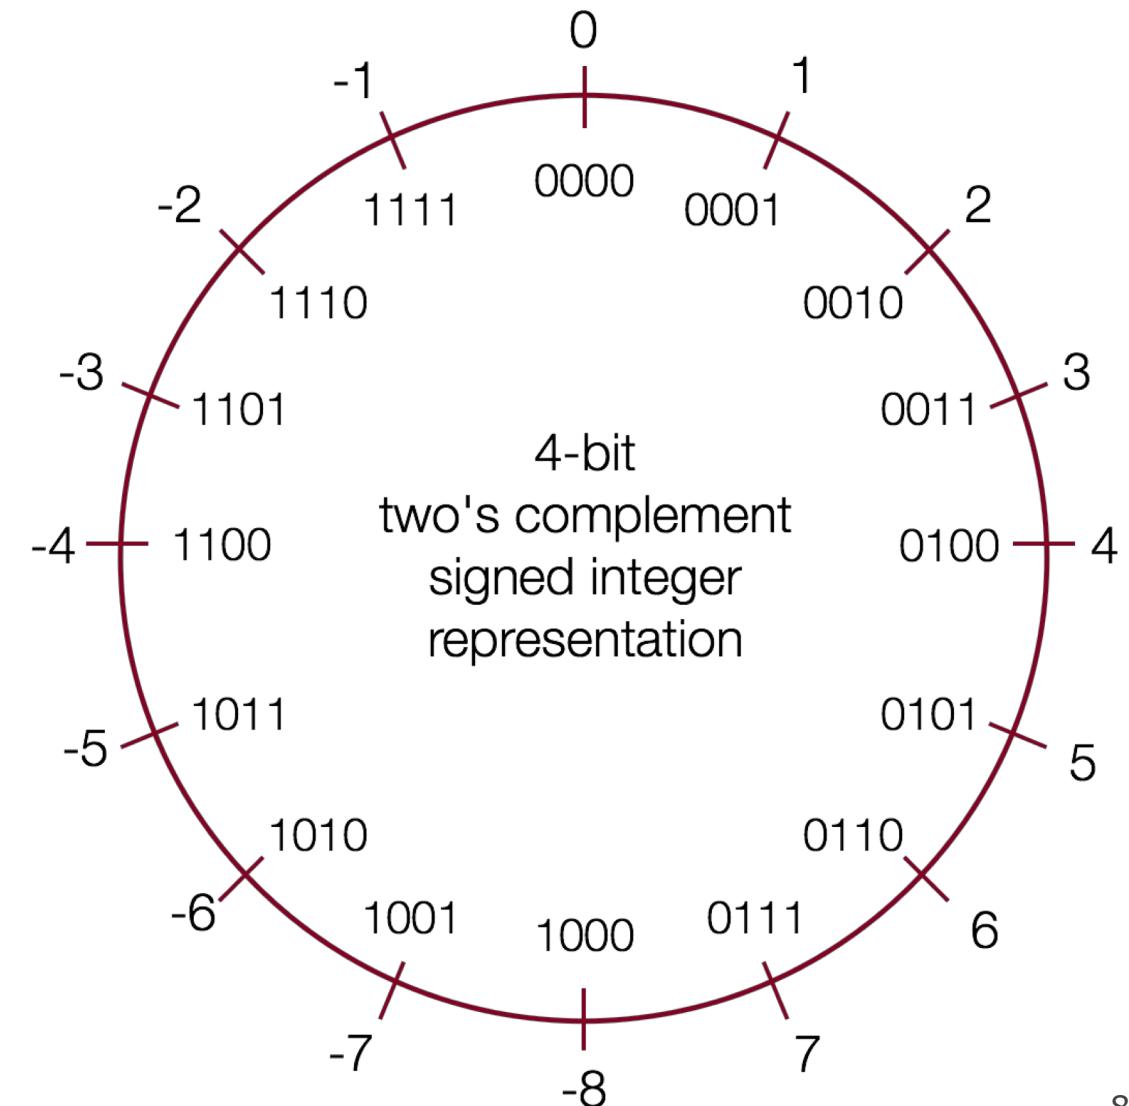
Myth

# Unsigned Integers



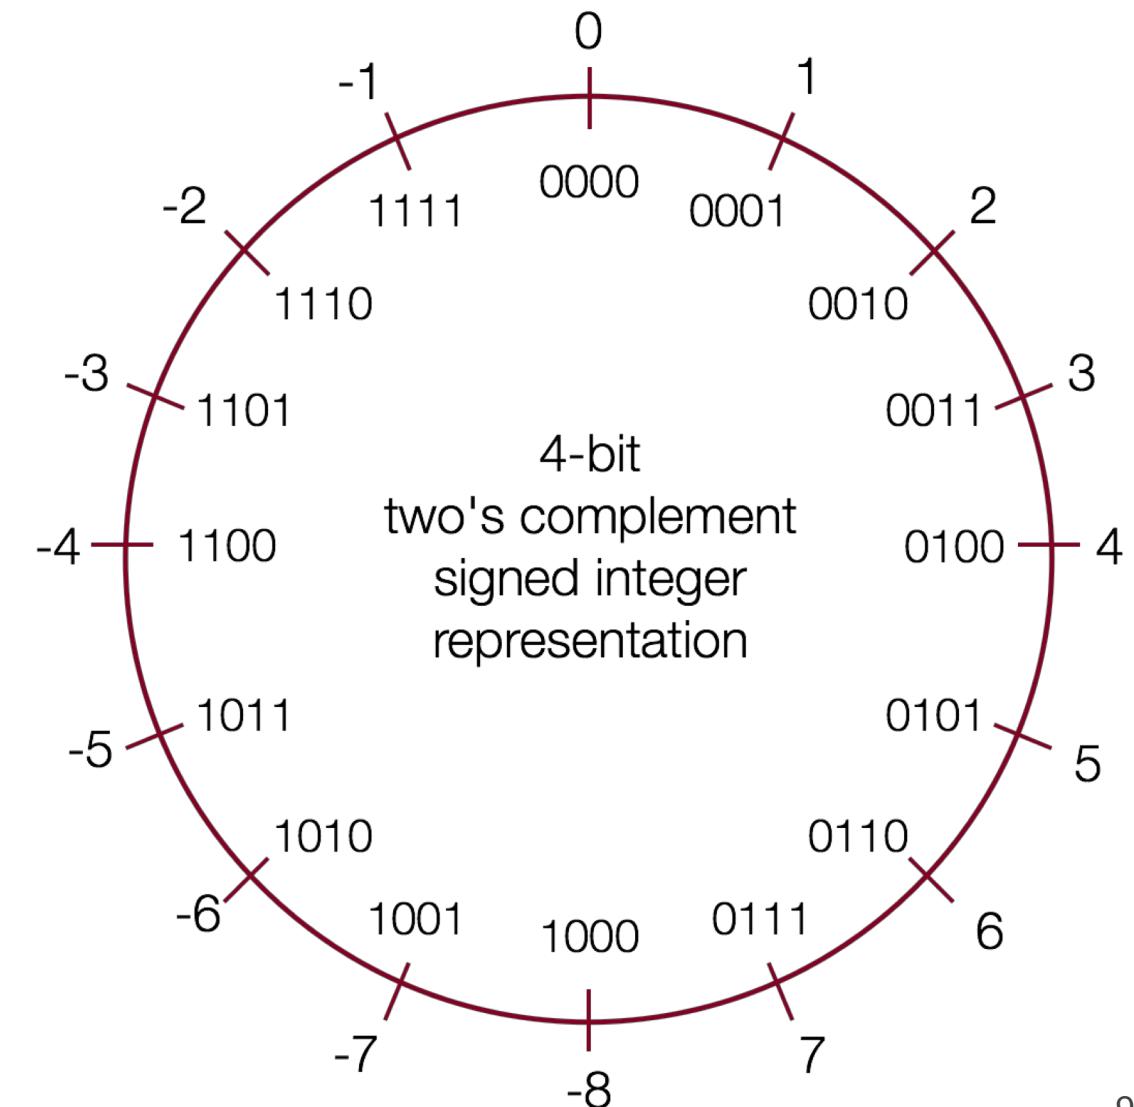
# Signed Integers: Two's Complement

- In **two's complement**, we represent a positive number as **itself**, and its negative equivalent as the **two's complement of itself**.
- The **two's complement** of a number is the binary digits inverted, plus 1.
- This works to convert from positive to negative, **and** back from negative to positive!



# Signed Integers: Two's Complement

- **Con:** more difficult to represent, and difficult to convert to/from decimal and between positive and negative.
- **Pro:** only 1 representation for 0!
- **Pro:** all bits are used to represent as many numbers as possible
- **Pro:** it turns out that the most significant bit *still indicates the sign* of a number.
- **Pro:** arithmetic is easy: we just add!



# Overflow and Underflow

- If you exceed the **maximum** value of your bit representation, you *wrap around* or *overflow* back to the **smallest** bit representation.

$$0b1111 + 0b1 = 0b0000$$

- If you go below the **minimum** value of your bit representation, you *wrap around* or *underflow* back to the **largest** bit representation.

$$0b0000 - 0b1 = 0b1111$$

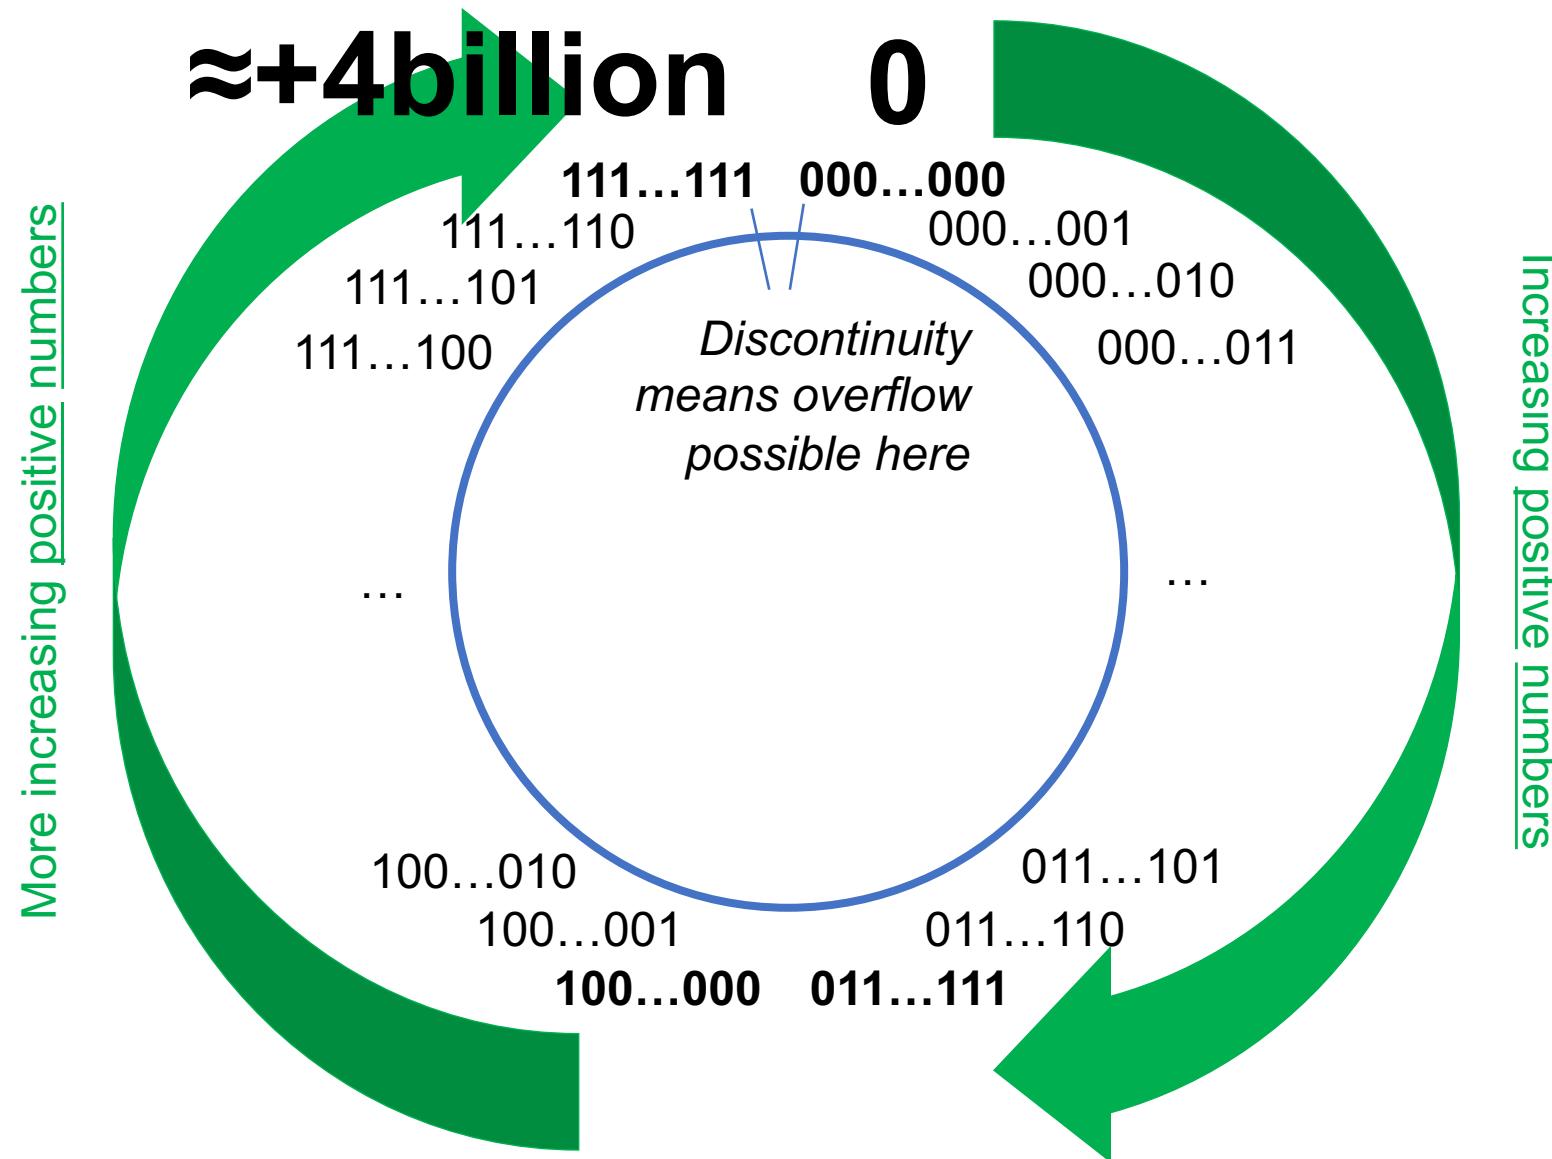
# Min and Max Integer Values

Type	Width (bytes)	Width (bits)	Min in hex (name)	Max in hex (name)
char	1	8	80 (CHAR_MIN)	7F (CHAR_MAX)
unsigned char	1	8	0	FF (UCHAR_MAX)
short	2	16	8000 (SHRT_MIN)	7FFF (SHRT_MAX)
unsigned short	2	16	0	FFFF (USHRT_MAX)
int	4	32	80000000 (INT_MIN)	7FFFFFFF (INT_MAX)
unsigned int	4	32	0	FFFFFFFF (UINT_MAX)
long	8	64	8000000000000000 (LONG_MIN)	7FFFFFFFFFFFFFFF (LONG_MAX)
unsigned long	8	64	0	FFFFFFFFFFFFFF (ULONG_MAX)

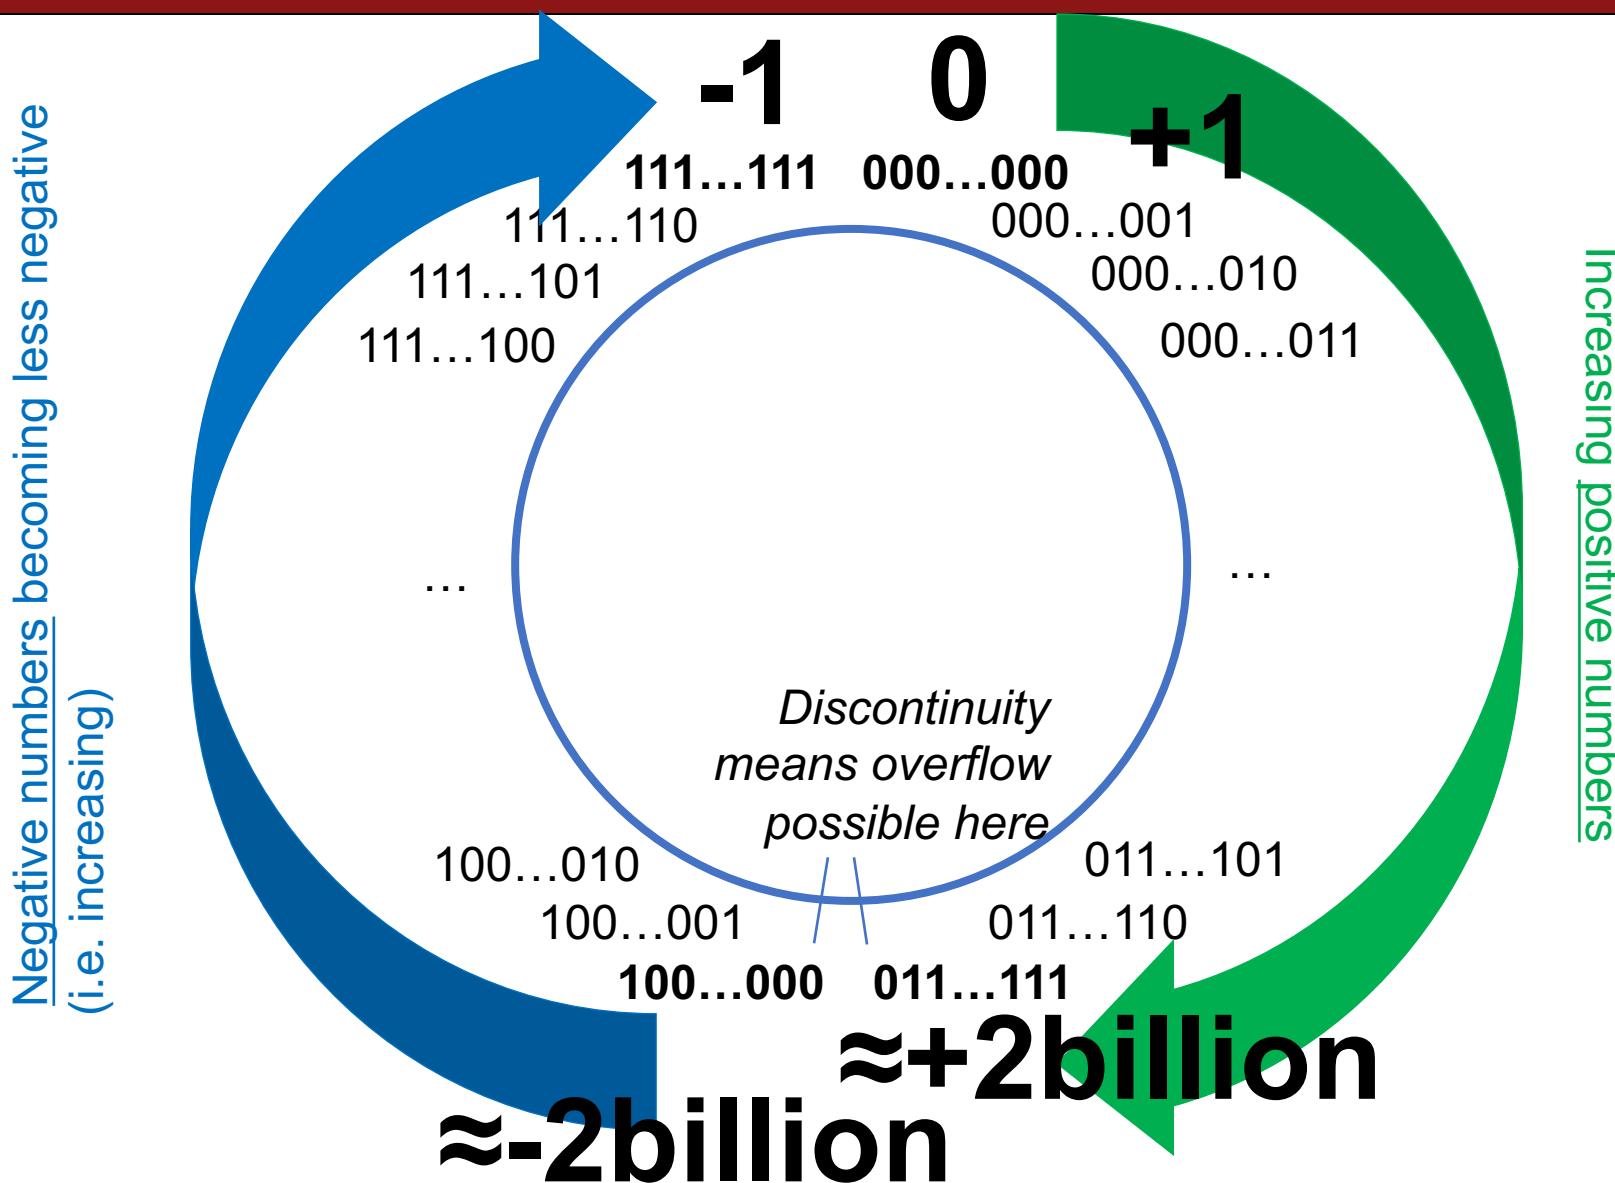
# Aside: ASCII

- ASCII is an encoding from common characters (letters, symbols, etc.) to bit representations (chars).
  - E.g. 'A' is 0x41
- Neat property: all uppercase letters, and all lowercase letters, are sequentially represented!
  - E.g. 'B' is 0x42

# Unsigned Integers



# Signed Numbers



# Casting

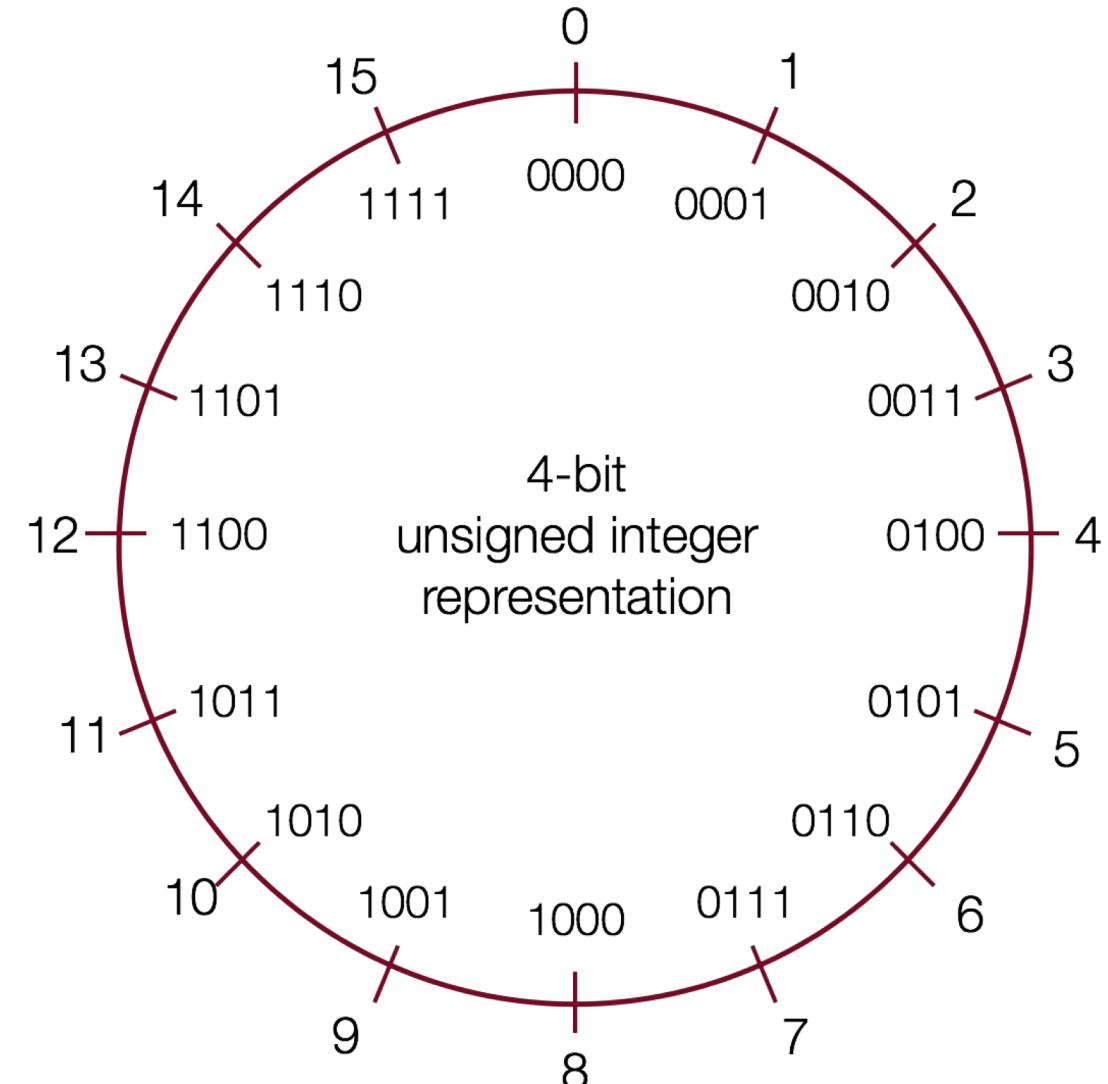
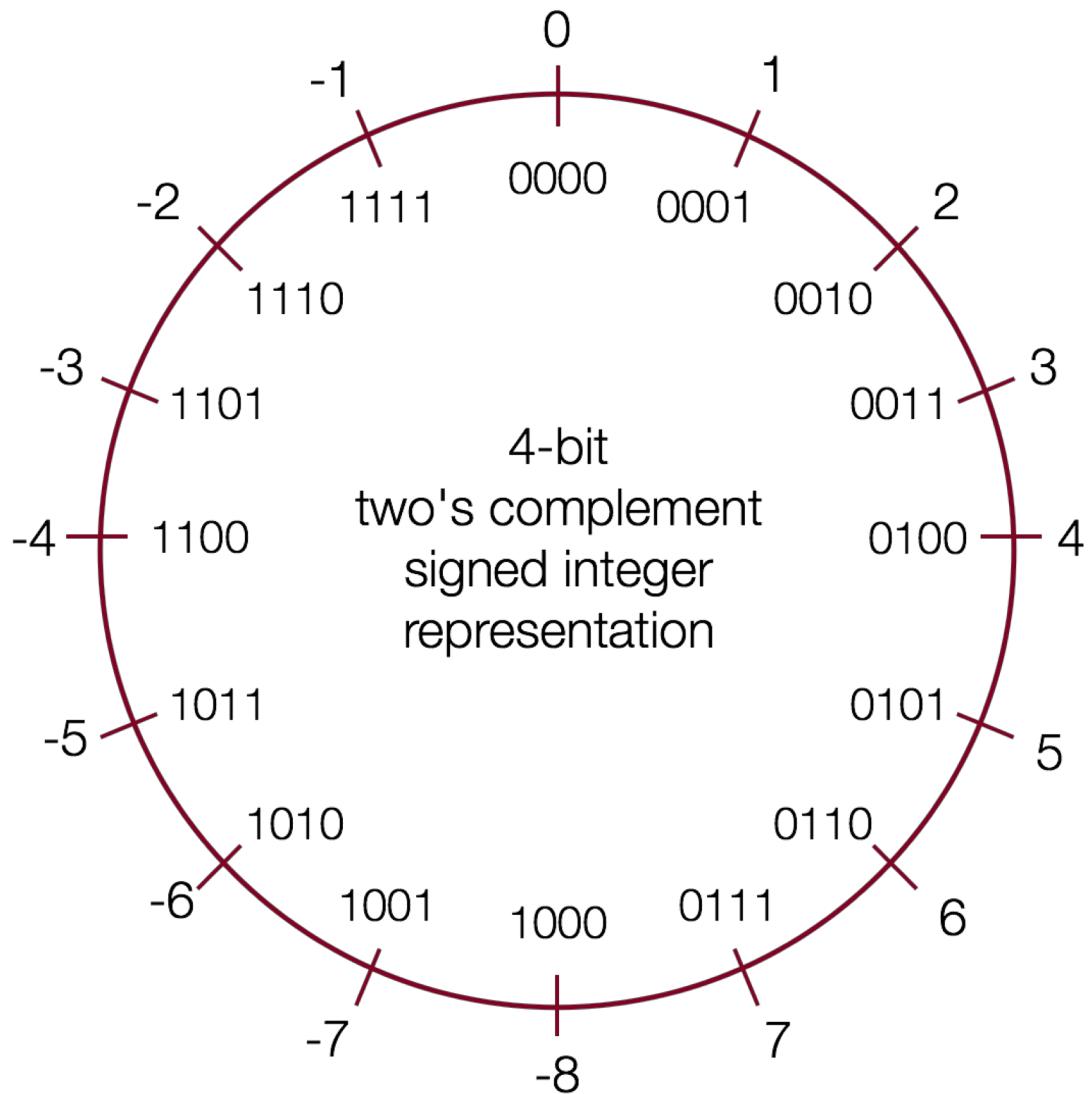
- What happens at the byte level when we cast between variable types? **The bytes remain the same! This means they may be interpreted differently depending on the type.**

```
int v = -12345;  
unsigned int uv = v;  
printf("v = %d, uv = %u\n", v, uv);
```

The bit representation for -12345 is **0b1100000111001**.

If we treat this binary representation as a positive number, it's *huge*!

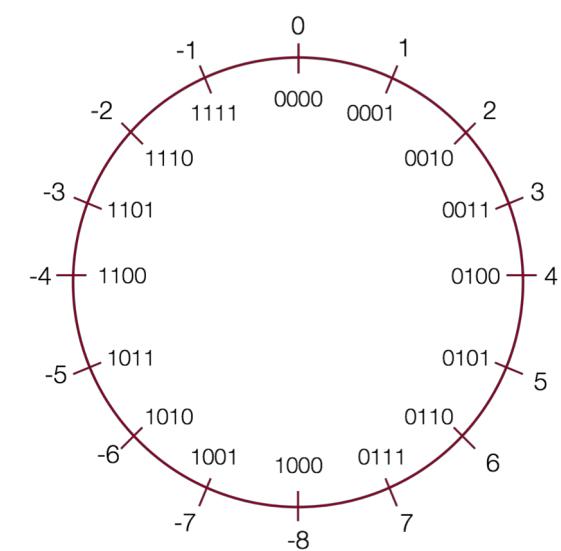
# Casting



# Comparisons Between Different Types

- Be careful when comparing signed and unsigned integers. C will implicitly cast the signed argument to unsigned, and then performs the operation assuming both numbers are non-negative.

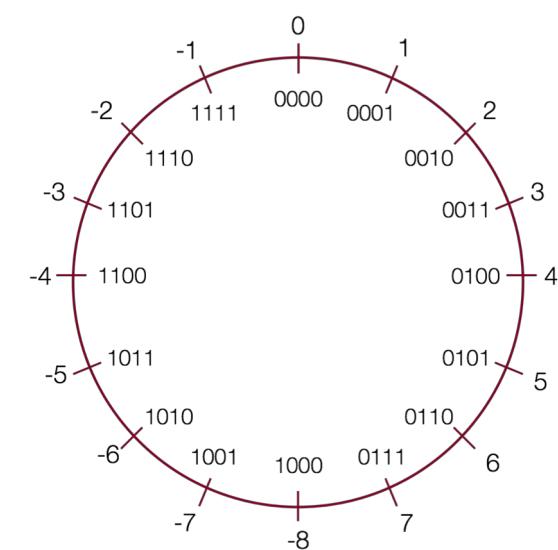
Expression	Type	Evaluation	Correct?
<code>0 == 0U</code>			
<code>-1 &lt; 0</code>			
<code>-1 &lt; 0U</code>			
<code>2147483647 &gt; -</code>			
<code>2147483647 - 1</code>			
<code>2147483647U &gt; -</code>			
<code>2147483647 - 1</code>			
<code>2147483647 &gt;</code>			
<code>(int)2147483648U</code>			
<code>-1 &gt; -2</code>			
<code>(unsigned)-1 &gt; -2</code>			



# Comparisons Between Different Types

- Be careful when comparing signed and unsigned integers. C will implicitly cast the signed argument to unsigned, and then performs the operation assuming both numbers are non-negative.

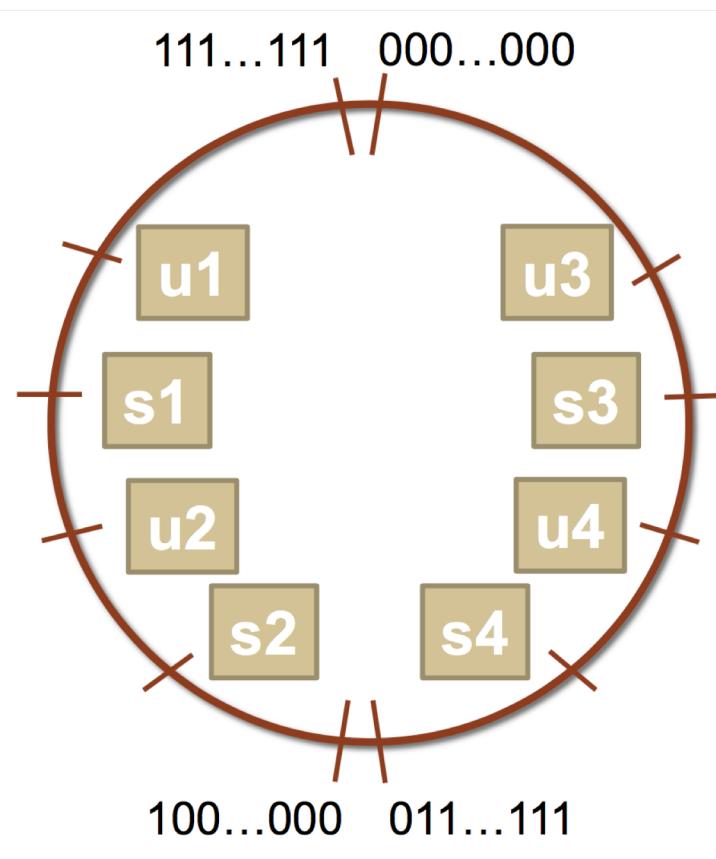
Expression	Type	Evaluation	Correct?
<code>0 == 0U</code>	Unsigned	1	yes
<code>-1 &lt; 0</code>	Signed	1	yes
<code>-1 &lt; 0U</code>	Unsigned	0	No!
<code>2147483647 &gt; -2147483647 - 1</code>	Signed	1	yes
<code>2147483647U &gt; -2147483647 - 1</code>	Unsigned	0	No!
<code>2147483647 &gt; (int)2147483648U</code>	Signed	1	No!
<code>-1 &gt; -2</code>	Signed	1	yes
<code>(unsigned)-1 &gt; -2</code>	Unsigned	1	yes



# Comparisons Between Different Types

Which many of the following statements are true? (assume that variables are set to values that place them in the spots shown)

- s3 > u3
- u2 > u4
- s2 > s4
- s1 > s2
- u1 > u2
- s1 > u3



# Comparisons Between Different Types

Which many of the following statements are true? (assume that variables are set to values that place them in the spots shown)

**s3 > u3**

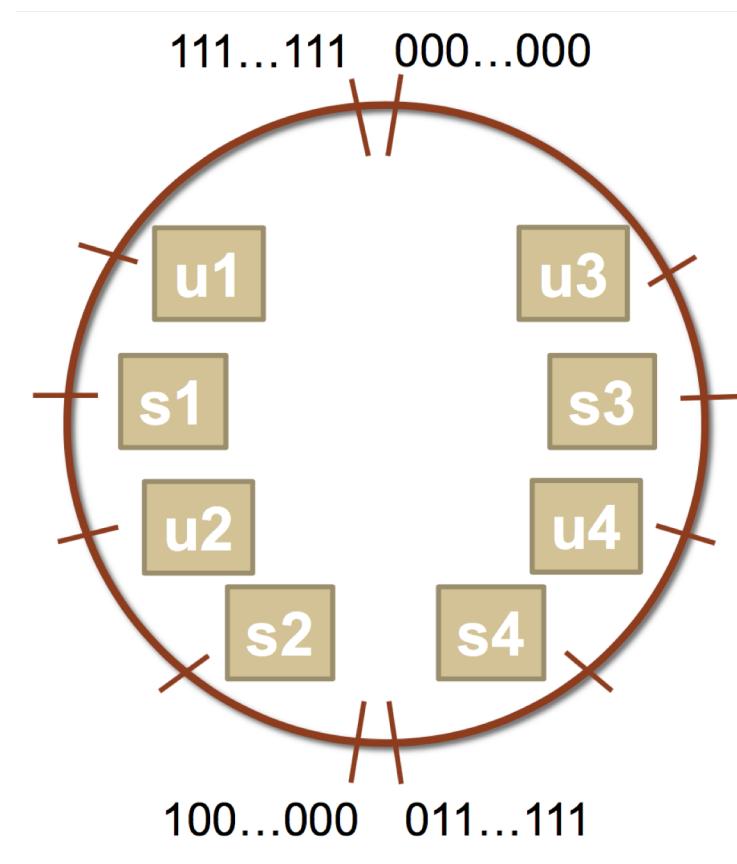
**u2 > u4**

**s2 > s4**

**s1 > s2**

**u1 > u2**

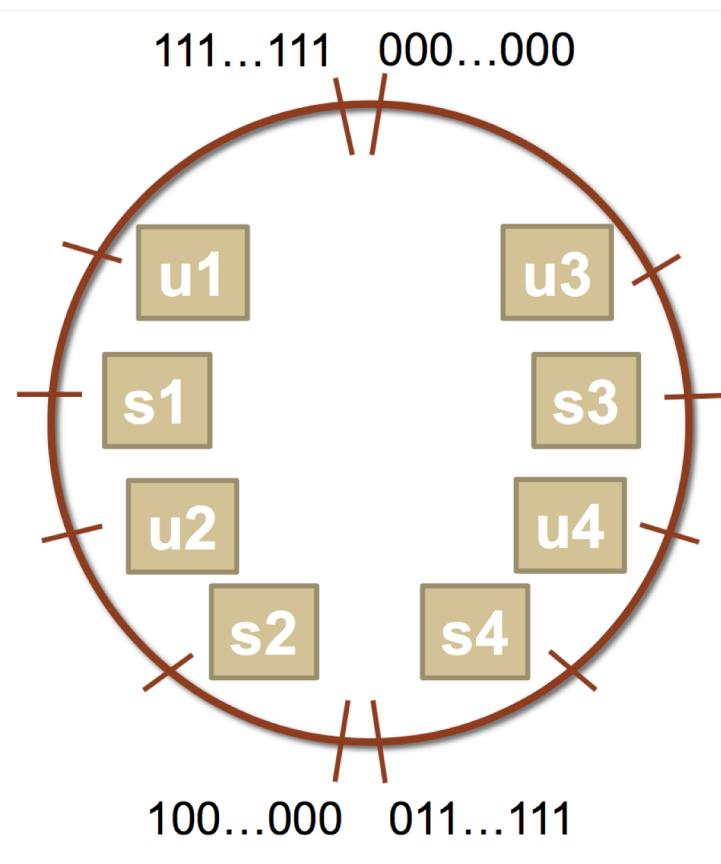
**s1 > u3**



# Comparisons Between Different Types

Which many of the following statements are true? (assume that variables are set to values that place them in the spots shown)

- s3 > u3 - true**
- u2 > u4 - true**
- s2 > s4 - false**
- s1 > s2 - true**
- u1 > u2 - true**
- s1 > u3 - true**



# Plan For Today

- **Recap:** Integer Representations
- **Truncating and Expanding**
- Bitwise Boolean Operators and Masks
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- Bit Shift Operators

# Expanding Bit Representations

- Sometimes, we want to convert between two integers of different sizes (e.g. short to int, or int to long).
- We might not be able to convert from a bigger data type to a smaller data type, but we do want to always be able to convert from a **smaller** data type to a **bigger** data type.
- For **unsigned** values, we can add *leading zeros* to the representation (“zero extension”)
- For **signed** values, we can *repeat the sign of the value* for new digits (“sign extension”)
- Note: when doing `<`, `>`, `<=`, `>=` comparison between different size types, it will *promote to the larger type*.

# Expanding Bit Representation

```
unsigned short s = 4;  
// short is a 16-bit format, so                                s = 0000 0000 0000 0100b  
  
unsigned int i = s;  
// conversion to 32-bit int, so i = 0000 0000 0000 0000 0000 0000 0000 0100b
```

# Expanding Bit Representation

```
short s = 4;  
// short is a 16-bit format, so s = 0000 0000 0000 0100b  
  
int i = s;  
// conversion to 32-bit int, so i = 0000 0000 0000 0000 0000 0000 0000 0100b  
  
— or —  
  
short s = -4;  
// short is a 16-bit format, so s = 1111 1111 1111 1100b  
  
int i = s;  
// conversion to 32-bit int, so i = 1111 1111 1111 1111 1111 1111 1111 1100b
```

# Truncating Bit Representation

If we want to **reduce** the bit size of a number, C *truncates* the representation and discards the *more significant bits*.

```
int x = 53191;  
short sx = x;  
int y = sx;
```

What happens here? Let's look at the bits in x (a 32-bit int), 53191:

**0000 0000 0000 0000 1100 1111 1100 0111**

When we cast x to a short, it only has 16-bits, and C *truncates* the number:

**1100 1111 1100 0111**

This is -12345! And when we cast sx back an int, we sign-extend the number.

**1111 1111 1111 1111 1100 1111 1100 0111 // still -12345**

# Truncating Bit Representation

If we want to **reduce** the bit size of a number, C *truncates* the representation and discards the *more significant bits*.

```
int x = -3;  
short sx = x;  
int y = sx;
```

What happens here? Let's look at the bits in x (a 32-bit int), -3:

1111 1111 1111 1111 1111 1111 1111 1101

When we cast x to a short, it only has 16-bits, and C *truncates* the number:

1111 1111 1111 1101

This is -3! **If the number does fit, it will convert fine.** y looks like this:

1111 1111 1111 1111 1111 1111 1111 1101 // still -3

# Truncating Bit Representation

If we want to **reduce** the bit size of a number, C *truncates* the representation and discards the *more significant bits*.

```
unsigned int x = 128000;  
unsigned short sx = x;  
unsigned int y = sx;
```

What happens here? Let's look at the bits in x (a 32-bit unsigned int), 128000:

0000 0000 0000 0001 1111 0100 0000 0000

When we cast x to a short, it only has 16-bits, and C *truncates* the number:

1111 0100 0000 0000

This is 62464! **Unsigned numbers can lose info too.** Here is what y looks like:

0000 0000 0000 0000 1111 0100 0000 0000 // still 62464

# The sizeof Operator

- **sizeof** takes a variable type as a parameter and returns the number of bytes that type uses.

```
printf("sizeof(char): %d\n", (int) sizeof(char));
printf("sizeof(short): %d\n", (int) sizeof(short));
printf("sizeof(int): %d\n", (int) sizeof(int));
printf("sizeof(unsigned int): %d\n", (int) sizeof(unsigned int));
printf("sizeof(long): %d\n", (int) sizeof(long));
printf("sizeof(long long): %d\n", (int) sizeof(long long));
printf("sizeof(size_t): %d\n", (int) sizeof(size_t));
printf("sizeof(void *): %d\n", (int) sizeof(void *));
```

```
$ ./sizeof
sizeof(char): 1
sizeof(short): 2
sizeof(int): 4
sizeof(unsigned int): 4
sizeof(long): 8
sizeof(long long): 8
sizeof(size_t): 8
sizeof(void *): 8
```

Type	Width in bytes	Width in bits
char	1	8
short	2	16
int	4	32
long	8	64
void *	8	64

# Plan For Today

- **Recap:** Integer Representations
- Truncating and Expanding
- **Bitwise Boolean Operators and Masks**
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- Bit Shift Operators

# Bitwise Operators

- You're already familiar with many operators in C:
  - **Arithmetic operators:** +, -, \*, /, %
  - **Comparison operators:** ==, !=, <, >, <=, >=
  - **Logical Operators:** &&, ||, !
- Today, we're introducing a new category of operators: **bitwise operators:**
  - &, |, ~, ^, <<, >>

# And (&)

AND is a binary operator. The AND of 2 bits is 1 if both bits are 1, and 0 otherwise. **Note:** this is different from Boolean AND (&&)!

**output = a & b;**

a	b	output
0	0	0
0	1	0
1	0	0
1	1	1

# Or (|)

OR is a binary operator. The OR of 2 bits is 1 if either (or both) bits is 1. **Note:** this is different from Boolean OR (||)!

**output = a | b;**

a	b	output
0	0	0
0	1	1
1	0	1
1	1	1

# Not ( $\sim$ )

NOT is a unary operator. The NOT of a bit is 1 if the bit is 0, or 1 otherwise.

**Note:** this is different from Boolean NOT (!)!

```
output = ~a;
```

a	output
0	1
1	0

# Exclusive Or (^)

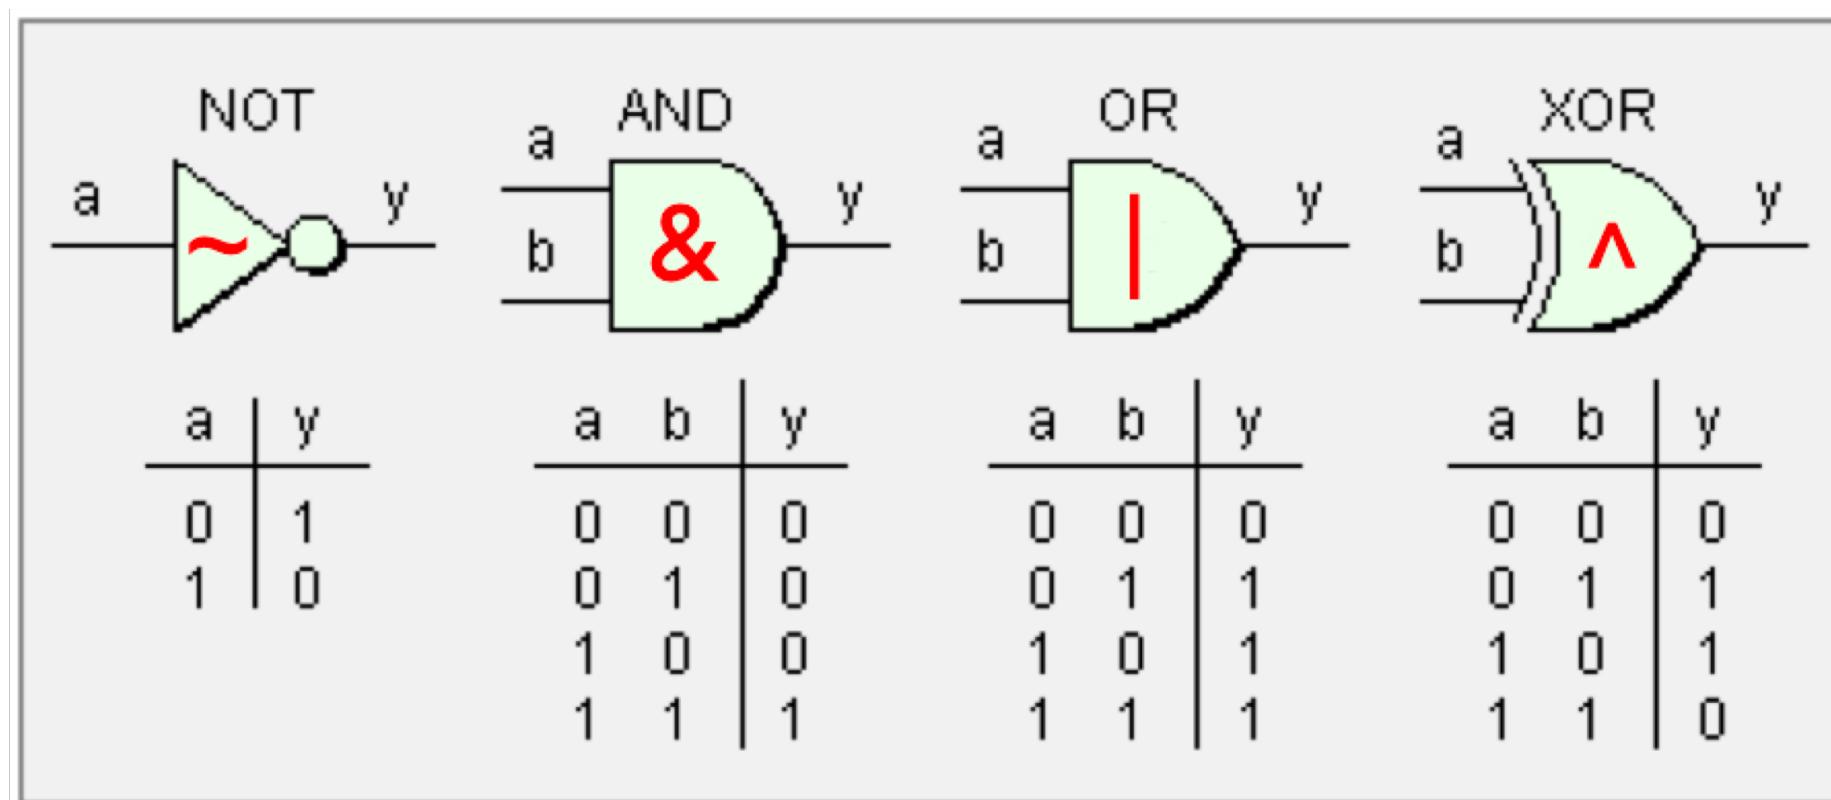
Exclusive Or (XOR) is a binary operator. The XOR of 2 bits is 1 if *exactly* one of the bits is 1, or 0 otherwise.

**output = a ^ b;**

a	b	output
0	0	0
0	1	1
1	0	1
1	1	0

# An Aside: Boolean Algebra

- These operators are not unique to computers; they are part of a general area called **Boolean Algebra**, and are called **Boolean Operators**. These are applicable in math, hardware, computers, and more!



# Operators on Multiple Bits

- When these Boolean operators are applied to numbers (multiple bits), the operator is applied to the corresponding bits in each number. For example:

AND

$$\begin{array}{r} 0110 \\ \& 1100 \\ \hline 0100 \end{array}$$

OR

$$\begin{array}{r} 0110 \\ | 1100 \\ \hline 1110 \end{array}$$

XOR

$$\begin{array}{r} 0110 \\ ^ 1100 \\ \hline 1010 \end{array}$$

NOT

$$\begin{array}{r} \sim 1100 \\ \hline 0011 \end{array}$$

# Bit Vectors and Sets

- We can use bit vectors (ordered collections of bits) to represent finite sets, and perform functions such as union, intersection, and complement.
- **Example:** we can represent current courses taken using a **char**.

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS110	CS101	CS106X	CS106B	CS106A

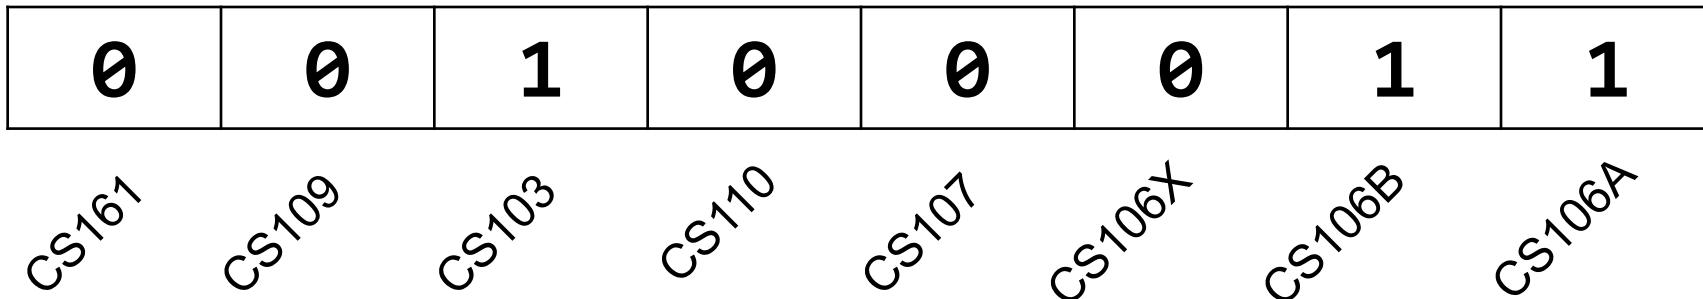
# Bit Vectors and Sets

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS110	CS107	CS106X	CS106B	CS106A

- How do we find the union of two sets of courses taken? Use OR:

$$\begin{array}{r} 00100011 \\ | \quad 01100001 \\ \hline 01100011 \end{array}$$

# Bit Vectors and Sets



- How do we find the intersection of two sets of courses taken? Use AND:

$$\begin{array}{r} 00100011 \\ \& 01100001 \\ \hline 00100001 \end{array}$$

# Bit Masking

- We will frequently want to manipulate or isolate out specific bits in a larger collection of bits. A **bitmask** is a constructed bit pattern that we can use, along with bit operators, to do this.
- **Example:** how do we update our bit vector to indicate we've taken CS107?

0	0	1	0	0	0	1	1
---	---	---	---	---	---	---	---

CS161

CS109

CS103

CS110

CS107

CS106X

CS106B

CS106A

00100011

| 00001000

-----

00101011

# Bit Masking

```
enum Classes {
    CS106A = 0x1,      /* 0000 0001 */
    CS106B = 0x2,      /* 0000 0010 */
    CS106X = 0x4,      /* 0000 0100 */
    CS107 = 0x8,       /* 0000 1000 */
    CS110 = 0x10,      /* 0001 0000 */
    CS103 = 0x20,      /* 0010 0000 */
    CS109 = 0x40,      /* 0100 0000 */
    CS161 = 0x80,      /* 1000 0000 */
};

char myClasses = ...;
myClasses = myClasses | CS107; // Add CS107
```

# Bit Masking

```
enum Classes {  
    CS106A = 0x1,      /* 0000 0001 */  
    CS106B = 0x2,      /* 0000 0010 */  
    CS106X = 0x4,      /* 0000 0100 */  
    CS107 = 0x8,       /* 0000 1000 */  
    CS110 = 0x10,      /* 0001 0000 */  
    CS103 = 0x20,      /* 0010 0000 */  
    CS109 = 0x40,      /* 0100 0000 */  
    CS161 = 0x80,      /* 1000 0000 */  
};
```

```
char myClasses = ...;  
myClasses |= CS107;      // Add CS107
```

# Bit Masking

- Example: how do we update our bit vector to indicate we've *not* taken CS103?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS110	CS101	CS106X	CS106B	CS106A

00100011  
& 11011111  
-----  
00000011

```
char myClasses = ...;  
myClasses = myClasses & ~CS103; // Remove CS103
```

# Bit Masking

- Example: how do we update our bit vector to indicate we've *not* taken CS103?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS110	CS101	CS106X	CS106B	CS106A

00100011  
& 11011111  
-----  
00000011

```
char myClasses = ...;  
myClasses &= ~CS103; // Remove CS103
```

# Bit Masking

- Example: how do we check if we've taken CS106B?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS110	CS101	CS106X	CS106B	CS106A

00100011  
& 00000010  
-----  
00000010

```
char myClasses = ...;  
if (myClasses & CS106B) {...  
    // taken CS106B!
```

# Bit Masking

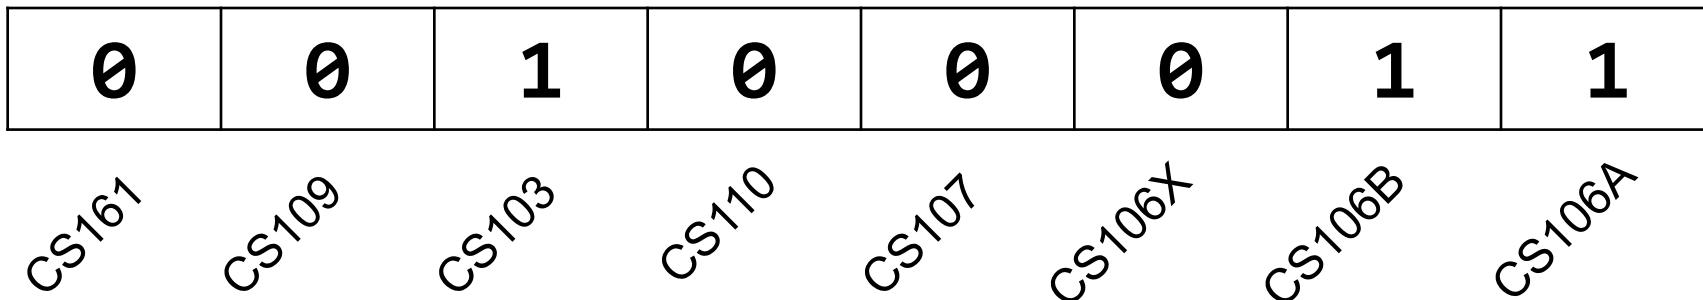
- Example: how do we check if we've *not* taken CS107?

0	0	1	0	0	0	1	1
CS161	CS109	CS103	CS110	CS107	CS106X	CS106B	CS106A

```
char myClasses = ...;
if ((myClasses & CS107) ^ CS107) {...  
    // not taken CS107!
```

# Bit Masking

- **Example:** how do we check if we've *not* taken CS107?



```
char myClasses = ...;  
if (!(myClasses & CS107)) {...  
    // not taken CS107!
```

# Demo: Bitmasks and GDB



# Bit Masking

- Bit masking is also useful for integer representations as well. For instance, we might want to check the value of the most-significant bit, or just one of the middle bytes.
- **Example:** If I have a 32-bit integer  $j$ , what operation should I perform if I want to get *just the lowest byte* in  $j$ ?

```
int j = ...;  
int k = j & 0xff;           // mask to get just lowest byte
```

# Practice: Bit Masking

- **Practice 1:** write an expression that, given a 32-bit integer  $j$ , sets its least-significant byte to all 1s, but preserves all other bytes.
- **Practice 2:** write an expression that, given a 32-bit integer  $j$ , flips (“complements”) all but the least-significant byte, and preserves all other bytes.

# Practice: Bit Masking

- **Practice 1:** write an expression that, given a 32-bit integer  $j$ , sets its least-significant byte to all 1s, but preserves all other bytes.

$j \mid 0xff$

- **Practice 2:** write an expression that, given a 32-bit integer  $j$ , flips (“complements”) all but the least-significant byte, and preserves all other bytes.

$j \wedge \sim 0xff$

# Plan For Today

- **Recap:** Integer Representations
- Truncating and Expanding
- Bitwise Boolean Operators and Masks
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- Bit Shift Operators

# Announcements

- Office Hours Updates – Nick’s OH location, Fri time, and future 1-time adjustments. *See office hours calendar!*
- Please send us any OAE letters or athletics conflicts as soon as possible. Thanks!
- Assignment 0 deadline tonight at 11:59PM PST
- Assignment 1 (Bit operations!) goes out tonight at Assignment 0 deadline
  - Saturated arithmetic
  - Game of Life
  - Unicode and UTF-8
- Lab 1 this week!

# Demo: Powers of 2



# Plan For Today

- **Recap:** Integer Representations
- Truncating and Expanding
- Bitwise Boolean Operators and Masks
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- **Bit Shift Operators**

# Left Shift (<<)

The LEFT SHIFT operator shifts a bit pattern a certain number of positions to the left. New lower order bits are filled in with 0s, and bits shifted off of the end are lost.

```
x << k;      // shifts x to the left by k bits
```

8-bit examples:

00110111 << 2 results in 11011100

01100011 << 4 results in 00110000

10010101 << 4 results in 01010000

# Right Shift (>>)

The RIGHT SHIFT operator shifts a bit pattern a certain number of positions to the right. Bits shifted off of the end are lost.

```
x >> k; // shifts x to the right by k bits
```

**Question:** how should we fill in new higher-order bits?

**Idea:** let's follow left-shift and fill with 0s.

```
short x = 2; // 0000 0000 0000 0010
x >> 1; // 0000 0000 0000 0001
printf("%d\n", x); // 1
```

# Right Shift (>>)

The RIGHT SHIFT operator shifts a bit pattern a certain number of positions to the right. Bits shifted off of the end are lost.

```
x >> k; // shifts x to the right by k bits
```

**Question:** how should we fill in new higher-order bits?

**Idea:** let's follow left-shift and fill with 0s.

```
short x = -2; // 1111 1111 1111 1110
x >> 1; // 0111 1111 1111 1111
printf("%d\n", x); // 32767!
```

# Right Shift (>>)

The RIGHT SHIFT operator shifts a bit pattern a certain number of positions to the right. Bits shifted off of the end are lost.

```
x >> k; // shifts x to the right by k bits
```

**Question:** how should we fill in new higher-order bits?

**Problem:** always filling with zeros means we may change the sign bit.

**Solution:** let's fill with the sign bit!

# Right Shift (>>)

The RIGHT SHIFT operator shifts a bit pattern a certain number of positions to the right. Bits shifted off of the end are lost.

```
x >> k; // shifts x to the right by k bits
```

**Question:** how should we fill in new higher-order bits?

**Solution:** let's fill with the sign bit!

```
short x = 2; // 0000 0000 0000 0010
x >> 1; // 0000 0000 0000 0001
printf("%d\n", x); // 1
```

# Right Shift (>>)

The RIGHT SHIFT operator shifts a bit pattern a certain number of positions to the right. Bits shifted off of the end are lost.

```
x >> k; // shifts x to the right by k bits
```

**Question:** how should we fill in new higher-order bits?

**Solution:** let's fill with the sign bit!

```
short x = -2; // 1111 1111 1111 1110
x >> 1; // 1111 1111 1111 1111
printf("%d\n", x); // -1!
```

# Right Shift (>>)

There are *two kinds* of right shifts, depending on the value and type you are shifting:

- **Logical Right Shift:** fill new high-order bits with 0s.
- **Arithmetic Right Shift:** fill new high-order bits with the most-significant bit.

*Unsigned numbers* are right-shifted using **Logical Right Shift**.

*Signed numbers* are right-shifted using **Arithmetic Right Shift**.

This way, the sign of the number (if applicable) is preserved!

# Shift Operation Pitfalls

1. *Technically*, the C standard does not precisely define whether a right shift for signed integers is logical or arithmetic. However, **almost all compilers/machines** use arithmetic, and you can most likely assume this.
2. Operator precedence can be tricky! For example:

**1<<2 + 3<<4** means **1 << (2+3) << 4** because addition and subtraction have higher precedence than shifts! Always use parentheses to be sure:

**(1<<2) + (3<<4)**

# Recap

- **Recap:** Integer Representations
- Truncating and Expanding
- Bitwise Boolean Operators and Masks
- **Demo 1:** Courses
- **Break:** Announcements
- **Demo 2:** Powers of 2
- Bit Shift Operators

**Next time:** C strings